

# Deep Learning PhD course

## HW#2

Muhammad Osama

June 14, 2019

## 1 MNIST

### 1.1

- Network Architecture: Dense 100 - Dense 10 (Taken similar to what I used in homework 1B)
- Optimizer: SGD with rate  $\eta = 0.25$
- Loss: Categorical cross entropy
- Metrics: Accuracy
- Total number of parameters =  $(784*100 + 100) + (100*10 + 10) = 79510$
- Epochs = 25, batch size = 100. Computation time  $\approx 50$  sec. In HW 1b, 6 number of epochs were used and same batch size but only the first epoch takes much longer than a minute. As far as accuracy is concerned, in HW1b at the end of 6 epochs the accuracy is around 97% whereas with this implementation, the accuracy is around 94% at the end of 6 epochs. But the final accuracy is similar.

### 1.2

Total number of parameters:

$$\underbrace{(3 * 3 * 1 + 1) * 8}_{1^{st} \text{ conv. layer}} + \underbrace{(3 * 3 * 8 + 1) * 16}_{2^{nd} \text{ conv. layer}} + \underbrace{(3 * 3 * 16 + 1) * 32}_{3^{rd} \text{ conv. layer}} + \underbrace{7 * 7 * 32 * 10 + 10}_{full \text{ layer}} = 21578$$




Figure 1

Layer (type)	Output Shape	Param #
conv2d_310 (Conv2D)	(None, 28, 28, 8)	80
activation_360 (Activation)	(None, 28, 28, 8)	0
max_pooling2d_148 (MaxPoolin	(None, 14, 14, 8)	0
conv2d_311 (Conv2D)	(None, 14, 14, 16)	1168
activation_361 (Activation)	(None, 14, 14, 16)	0
max_pooling2d_149 (MaxPoolin	(None, 7, 7, 16)	0
conv2d_312 (Conv2D)	(None, 7, 7, 32)	4640
activation_362 (Activation)	(None, 7, 7, 32)	0
flatten_6 (Flatten)	(None, 1568)	0
dense_8 (Dense)	(None, 10)	15690
activation_363 (Activation)	(None, 10)	0
Total params: 21,578		
Trainable params: 21,578		
Non-trainable params: 0		

Figure 2

The number of parameters are much less as compared to the network in section 1.1.

### 1.3

Activation function: tanh

Optimizer: SGD with rate  $\eta = 0.25$

The only difference that can be seen from figure 4 output is that when the

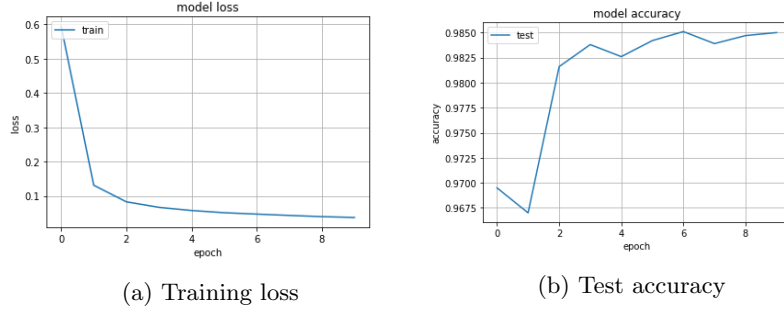


Figure 3

```
-- Activation layer before the pooling layer
Train on 60000 samples, validate on 10000 samples
Epoch 1/4
60000/60000 [=====] - 28s 463us/step - loss: 0.2949 - acc: 0.9070 - val_loss: 0.0936 - val_acc: 0.9704
Epoch 2/4
60000/60000 [=====] - 25s 416us/step - loss: 0.0859 - acc: 0.9731 - val_loss: 0.0612 - val_acc: 0.9793
Epoch 3/4
60000/60000 [=====] - 25s 419us/step - loss: 0.0613 - acc: 0.9808 - val_loss: 0.0404 - val_acc: 0.9861
Epoch 4/4
60000/60000 [=====] - 25s 414us/step - loss: 0.0503 - acc: 0.9847 - val_loss: 0.0397 - val_acc: 0.9866

--Activation layer after the pooling layer
Train on 60000 samples, validate on 10000 samples
Epoch 1/4
60000/60000 [=====] - 28s 469us/step - loss: 0.2947 - acc: 0.9071 - val_loss: 0.0936 - val_acc: 0.9705
Epoch 2/4
60000/60000 [=====] - 24s 394us/step - loss: 0.0859 - acc: 0.9731 - val_loss: 0.0611 - val_acc: 0.9792
Epoch 3/4
60000/60000 [=====] - 24s 397us/step - loss: 0.0612 - acc: 0.9808 - val_loss: 0.0404 - val_acc: 0.9860
Epoch 4/4
60000/60000 [=====] - 24s 396us/step - loss: 0.0502 - acc: 0.9847 - val_loss: 0.0397 - val_acc: 0.9865
```

Figure 4

activation layer is after the pooling layer the computational time is slightly less. This is possible because the activation is computed for a smaller image in the later case and hence some time is saved.

## 1.4

Activation function: relu

Optimizer: Adam with default values of rate  $\eta = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ ,  $\epsilon = 10^{-7}$ .




Figure 5

## 1.5

### 1.5.1

Activation changed from 'relu' to 'tanh'. Figure 6.




Figure 6

### 1.5.2

Activation: relu

Drop 50% of the weights. Figure 7

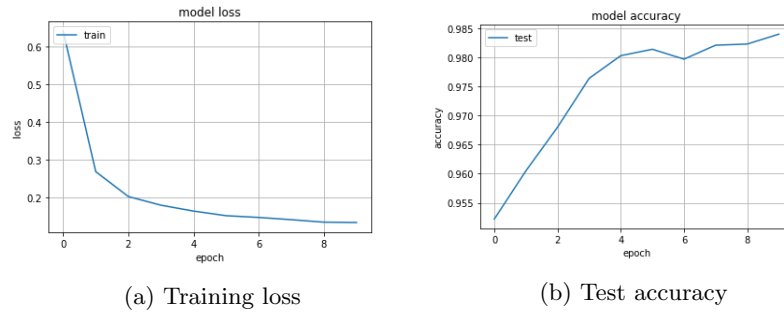


Figure 7

### 1.5.3

Optimizer:SGD with rate  $\eta = 0.001$  and momentum  $\alpha = 0.8$ . Figure 8

Confusion matrix for the best model which was obtained with the activation function 'tanh' and is shown in figure 9. A few examples of misclassified images are shown in figure 10.

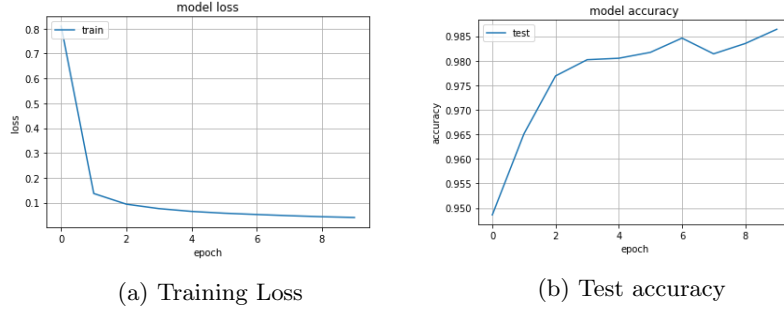


Figure 8




Figure 9: Confusion matrix for MNIST

## 2 WARWICK

### 2.1

The model summary for the model used is given in figure 11. The activation used was *relu*. The optimizer is *Adam* with it default values of learning rate  $\eta = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-7}$ .

The loss function was modified to handle pixel wise class maps by defining the a new function. A function to calculate the Dice coefficient was also defined as shown in figure 12.

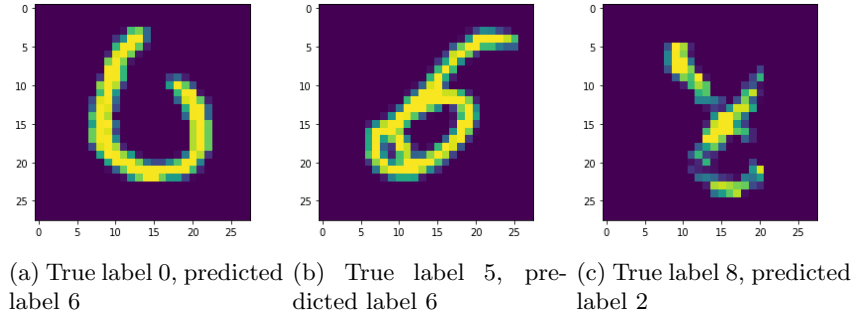


Figure 10

Layer (type)	Output Shape	Param #
conv2d_275 (Conv2D)	(None, 128, 128, 8)	224
activation_308 (Activation)	(None, 128, 128, 8)	0
max_pooling2d_130 (MaxPoolin	(None, 64, 64, 8)	0
conv2d_276 (Conv2D)	(None, 64, 64, 16)	1168
activation_309 (Activation)	(None, 64, 64, 16)	0
max_pooling2d_131 (MaxPoolin	(None, 32, 32, 16)	0
conv2d_277 (Conv2D)	(None, 32, 32, 32)	4640
activation_310 (Activation)	(None, 32, 32, 32)	0
conv2d_transpose_122 (Conv2D	(None, 64, 64, 16)	2064
activation_311 (Activation)	(None, 64, 64, 16)	0
conv2d_transpose_123 (Conv2D	(None, 128, 128, 8)	520
activation_312 (Activation)	(None, 128, 128, 8)	0
conv2d_278 (Conv2D)	(None, 128, 128, 1)	9
activation_313 (Activation)	(None, 128, 128, 1)	0
Total params: 8,625		
Trainable params: 8,625		
Non-trainable params: 0		

Figure 11

### 2.1.1

The Dice coefficient goes above 0.7 in approximately 200 iterations as shown in figure 13.

Figures 14 and 15 show two examples of test images which had low Dice coefficient. I think the reason that these test images have low Dice coefficient is because in most of the training images the glands are distributed over the space. But in these specific cases in the test data, the gland area is concentrated. This is also suggested by the predicted class map since it sort of gives a spatially

```

from keras import backend as K
from keras.activations import softmax

_EPSILON = K.epsilon()
def pixel_wise_loss(y_true, y_pred):
    y_pred = K.clip(y_pred, _EPSILON, 1.0-_EPSILON)

    loss = y_true*K.log(y_pred) + (1-y_true) * K.log(1-y_pred)
    #loss = K.sum(loss,axis=1)
    #loss = K.sum(loss,axis=1)
    return -loss

def dice_coef(y_true, y_pred):
    """
    Dice = (2*|X & Y|)/ (|X|+ |Y|)
          = 2*sum(|A*B|)/(sum(A^2)+sum(B^2))
    """
    y_true = K.flatten(y_true)
    y_pred = K.flatten(y_pred)
    intersection = K.sum(y_true * y_pred)
    return (2 * intersection) / (K.sum(y_true) + K.sum(y_pred))

```

Figure 12




Figure 13

distributed class map.




Figure 14: (a) Actual image (b) True class map (c) Predicted class map




Figure 15: (a) Actual image (b) True class map (c) Predicted class map

## 2.2

The Dice coefficient for the three different approaches is shown in figure 16.

(a) Data augmentation: Horizontally flipped the images and class maps of the training data and augmented to the original. The training data size doubles and its advantage can be clearly seen in figure 16a. The dice coefficient on the test data becomes greater than 0.7 in only 25 epochs where as in figure 13 it took almost 200 epochs.

(b) Dropping out 30% of the units in each convolution layer. The number of parameters in the original network were only approximately 8000 therefore when we drop out the units that the Dice coefficient increases slowly.

(c) Weight regularization using a  $\lambda = 10^{-4}$




Figure 16: (a) Data augmentation (b) Dropout(0.3) (c) Regularization



## Appendix

### Task1.1

```
import load_mnist as data
import numpy as np
from keras.models import Sequential
from keras.layers import Dense,Dropout
from keras import optimizers
import matplotlib.pyplot as plt

#load data
#X_train , Y_train , X_test , Y_test = data.
    load_mnist();

#reshape
#X_train = X_train.reshape(60000,28,28,1)
#X_test = X_test.reshape(10000,28,28,1)

# fix random seed for reproducibility
seed = 7
np.random.seed(seed)

#create model
model = Sequential()
#add model layers
model.add(Dense(100, activation='sigmoid
    ',input_dim=784))
model.add(Dense(10, activation='softmax')
    )

#optimizer
sgd = optimizers.SGD(lr=0.25, momentum
    =0.0, decay=0.0, nesterov=False);

#compile model using accuracy to measure
    model performance
model.compile(optimizer=sgd , loss='
    categorical_crossentropy ', metrics=['
    accuracy '])

#train the model
history = model.fit(X_train , Y_train ,
    validation_data = (X_test ,Y_test),
    epochs=25,batch_size = 100)
```

```

# list all data in history
print(history.history.keys())

# summarize history for accuracy
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.grid()
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper
left')
plt.show()

# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.grid()
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper
left')
plt.show()

```

## Task1.2

```
import load_mnist as data
import numpy as np
from keras.models import Sequential
from keras.layers import Conv2D,
    Activation, MaxPooling2D,
    ZeroPadding2D, Flatten, Dense
from keras import optimizers
import matplotlib.pyplot as plt
import funcs_segmentation as fs

#load data
#X_train, Y_train, X_test, Y_test = data.
    load_mnist();

#reshape
#X_train = X_train.reshape(60000,28,28,1)
#X_test = X_test.reshape(10000,28,28,1)

# fix random seed for reproducibility
seed = 7
np.random.seed(seed)

#create model
model = Sequential()
#add model layers
#1
#model.add(ZeroPadding2D(padding=1,
    data_format = 'channels_last'))
model.add(Conv2D(8, kernel_size=(3,3),
    strides=1, padding='same', data_format='
    channels_last', use_bias='True',
    kernel_initializer='random_normal',
    bias_initializer='zeros', input_shape
    =(28,28,1)))
#2
model.add(Activation('relu'))
#3
model.add(MaxPooling2D(pool_size=(2, 2),
    strides=2, padding='valid',
    data_format='channels_last'))
#4
#model.add(ZeroPadding2D(padding=1,
    data_format = 'channels_last'))
model.add(Conv2D(16, kernel_size=(3,3),
```

```

        strides=1,padding='same',data_format='
        channels_last ',use_bias='True',
        kernel_initializer='random_normal',
        bias_initializer='zeros'))
#5
model.add(Activation('relu'))
#6
model.add(MaxPooling2D(pool_size=(2, 2),
        strides=2, padding='valid ',
        data_format='channels_last'))
#7
#model.add(ZeroPadding2D(padding=1,
        data_format = 'channels_last'))
model.add(Conv2D(32,kernel_size=(3,3),
        strides=1,padding='same',data_format='
        channels_last ',use_bias='True',
        kernel_initializer='random_normal',
        bias_initializer='zeros'))
#8
model.add(Activation('relu'))
#
model.add(Flatten(data_format='
        channels_last'))
#9
model.add(Dense(10,use_bias='True',
        kernel_initializer='random_normal',
        bias_initializer='zeros'))
#10
model.add(Activation('softmax'))

#optimizer
sgd = optimizers.SGD(lr=0.25, momentum
        =0.0, decay=0.0, nesterov=False);

#compile model using accuracy to measure
        model performance
model.compile(optimizer=sgd , loss='
        categorical_crossentropy', metrics=['
        accuracy'])

#train the model
history = model.fit(X_train, Y_train,
        validation_data = (X_test, Y_test),
        epochs=4,batch_size = 100)

#get accuracy on test data

```

```

#model.evaluate(X_test,Y_test)

# list all data in history
#print(history.history.keys())

# summarize history for accuracy
#plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.grid()
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['test'], loc='upper left')
plt.show()

# summarize history for loss
plt.plot(history.history['loss'])
#plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.grid()
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()

```

For tasks 1.3 to 1.5 small changes were done to the file for task1.2 as mentioned in the report. Therefore it would just be redundant to paste their code here.

#### Task 2-Warwick

```
import load_warwick_augmented as data
#import numpy as np
from keras.models import Sequential
from keras.layers import Conv2D,
    Activation, MaxPooling2D,
    Conv2DTranspose#, Dropout
from keras import optimizers
#from keras.regularizers import l2
import funcs_segmentation as fs
import matplotlib.pyplot as plt

#load data
X_train, Y_train, X_test, Y_test = data.
    load_warwick();

# fix random seed for reproducibility
#seed = 7
#np.random.seed(seed)

#create model
model = Sequential()
#add model layers
#####
model.add(Conv2D(8, kernel_size=(3,3),
    strides=1,padding='same', data_format='
    channels_last', input_shape=(128,128,3)
))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2),
    strides=2, padding='valid',
    data_format='channels_last'))

#####
model.add(Conv2D(16, kernel_size=(3,3),
    strides=1,padding='same', data_format='
    channels_last'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2),
    strides=2, padding='valid',
    data_format='channels_last'))
```

```

#####
model.add(Conv2D(32, kernel_size=(3,3),
                strides=1, padding='same', data_format
                ='channels_last'))
model.add(Activation('relu'))

#####
model.add(Conv2DTranspose(16, kernel_size
                        =(2,2), strides=(2, 2), padding='same
                        ', data_format='channels_last'));
model.add(Activation('relu'))

model.add(Conv2DTranspose(8, kernel_size
                        =(2,2), strides=(2, 2), padding='same
                        ', data_format='channels_last'));
model.add(Activation('relu'))

model.add(Conv2D(1, kernel_size = (1,1)))

model.add(Activation('sigmoid'))

#optimizer
sgd = optimizers.SGD(lr=0.001, momentum
                    =0.0, decay=0.0, nesterov=False);

#compile model using accuracy to measure
model performanc
model.compile(optimizer='adam', loss=fs.
            pixel_wise_loss, metrics=[fs.dice_coef
            ])

#train the model
history = model.fit(X_train, Y_train,
                    validation_data = (X_test, Y_test),
                    epochs=250, batch_size = 5)

#get accuracy on test data
#model.evaluate(X_test, Y_test)

# list all data in history
#print(history.history.keys())

# summarize history for accuracy
plt.plot(history.history['dice_coef'])
plt.plot(history.history['val_dice_coef'])

```

```

    '])
plt.grid()
plt.title('Dice Coefficient')
#plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train','test'], loc='upper
           left')
plt.show()

# summarize history for loss
#plt.plot(history.history['loss'])
#plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.grid()
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()

```



## Supporting functions for task2

```
from keras import backend as K
from keras.activations import softmax

_EPSILON = K.epsilon()
def pixel_wise_loss(y_true, y_pred):
    y_pred = K.clip(y_pred, _EPSILON, 1.0 -
                     _EPSILON)

    loss = y_true*K.log(y_pred) + (1-y_true)
           * K.log(1-y_pred)
    #loss = K.sum(loss, axis=1)
    #loss = K.sum(loss, axis=1)
    return -loss

def dice_coef(y_true, y_pred):
    """
    Dice = (2*|X & Y|) / (|X|+ |Y|)
    = 2*sum(|A*B|)/(sum(A^2)+sum(B^2))
    """
    y_true = K.flatten(y_true)
    y_pred = K.flatten(y_pred)
    intersection = K.sum(y_true * y_pred)
    return (2 * intersection) / (K.sum(y_true
                                         ) + K.sum(y_pred))

def my_cross_entrp(y_true, y_pred):
    y_pred = K.clip(y_pred, _EPSILON, 1.0 -
                     _EPSILON)

    loss = K.sum(y_true*K.log(y_pred), axis=1)

    return -loss

def dice_loss(y_true, y_pred):

    dloss = 1-dice_coef(y_true, y_pred)

    return dloss
```

```
def softMaxAxis1(x):  
    return softmax(x,axis=2)
```

## Data Augmentation

```
import numpy as np
import matplotlib.pyplot as misc

def load_warwick():
    # Loads the MNIST dataset from png images

    NUM_TEST_IMAGES = 60
    # create list of image objects
    test_images = []
    test_class_maps = []

    for label in np.arange(1, NUM_TEST_IMAGES
        +1):
        if label < 10:
            image_path= "WARWICK/WARWICK/Test/image_"
                + str(0)+str(label) + ".png"
            class_map_path= "WARWICK/WARWICK/Test/"
                label_ + str(0)+str(label) + ".png"
        else:
            image_path= "WARWICK/WARWICK/Test/image_"
                + str(label) + ".png"
            class_map_path= "WARWICK/WARWICK/Test/"
                label_ + str(label) + ".png"

        image = misc.imread(image_path)
        class_map = misc.imread(class_map_path)
        test_images.append(image)
        test_class_maps.append(class_map)

    # create list of image objects
    NUM_TRN_IMAGES = 85

    train_images = []
    train_class_maps = []

    for label in np.arange(1, NUM_TRN_IMAGES
        +1):
        if label < 10:
            image_path= "WARWICK/WARWICK/Train/image_"
                + str(0)+str(label) + ".png"
            class_map_path= "WARWICK/WARWICK/Train/"
                label_ + str(0)+str(label) + ".png"
        else:
            image_path= "WARWICK/WARWICK/Train/image_"
```

```

    " + str(label) + ".png"
class_map_path= "WARWICK/WARWICK/Train/
    label_" + str(label) + ".png"

image = misc.imread(image_path)
class_map = misc.imread(class_map_path)
train_images.append(image)
train_class_maps.append(class_map)
train_images.append(np.fliplr(image))
train_class_maps.append(np.fliplr(
    class_map))

X_train= np.array(train_images).reshape
    (85*2,128,128,3)
Y_train= np.array(train_class_maps).
    reshape(85*2,128,128,1)
X_test= np.array(test_images).reshape
    (60,128,128,3)
Y_test= np.array(test_class_maps).reshape
    (60,128,128,1)

return X_train, Y_train, X_test, Y_test

```