

Deep Learning PhD course

HW#3

Muhammad Osama

June 14, 2019

1 Ex1

Assuming $\mathbf{x}^{(i)}$ is d dimensional and is either 0 or 1, we use the following data model conditioned on the latent variable \mathbf{z}

$$p_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}|\mathbf{z}) = \prod_{k=1}^d \text{Bernoulli}(\theta_k) \quad (1)$$

where $\boldsymbol{\theta} = \{\theta_1, \dots, \theta_d\}$. For *MNIST*, $d = 784$. Hence given \mathbf{z} , we take every pixel to be a Bernoulli random variable. Ideally, the pixel values lie in the interval $[0, 1]$. However, for simplicity we use the above approximation.

2 Ex2

The reason we cannot work with eq. (5) directly is because it is difficult to evaluate the expectation $\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})}[\log p_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}|\mathbf{z})]$ analytically.

3 Ex 3

$$q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}) = \mathcal{N}(\boldsymbol{\mu}^{(i)}, \text{diag}(\boldsymbol{\sigma}^{2(i)})) \quad p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$$

Let $\Sigma = \text{diag}(\boldsymbol{\sigma}^{2(i)})$, hence the KL divergence is

$$\begin{aligned} \Delta_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})||p(\mathbf{z})) &= \mathbb{E}_{q_{\phi}(\cdot)} \left[-\frac{(\mathbf{z} - \boldsymbol{\mu}^{(i)})^{\top} \Sigma^{-1} (\mathbf{z} - \boldsymbol{\mu}^{(i)})}{2} \right] + \\ &\quad \mathbb{E}_{q_{\phi}(\cdot)} \left[\frac{\mathbf{z}^{\top} \mathbf{z}}{2} \right] - \frac{1}{2} \sum_{j=1}^K \log \sigma_j^{2(i)}, \end{aligned} \quad (2)$$

where

$$\mathbb{E}_{q_\phi(\cdot)} \left[-\frac{(\mathbf{z} - \boldsymbol{\mu}^{(i)})^\top \Sigma^{-1} (\mathbf{z} - \boldsymbol{\mu}^{(i)})}{2} \right] = \frac{1}{2} \sum_{j=1}^K (1),$$

$$\mathbb{E}_{q_\phi(\cdot)} \left[\frac{\mathbf{z}^\top \mathbf{z}}{2} \right] = \sum_{j=1}^K \sigma_j^{2(i)} + \mu_j^{2(i)}.$$

Hence,

$$\Delta_{KL}(q_\phi(\mathbf{z}|\mathbf{x}^{(i)})||p(\mathbf{z})) = -\frac{1}{2} \left[\sum_{j=1}^K 1 + \log \sigma_j^{2(i)} - \sigma_j^{2(i)} - \mu_j^{2(i)} \right] \quad (3)$$

Moreover,

$$g_\phi(\boldsymbol{\epsilon}^{(i,l)}, \mathbf{x}^{(i)}) = \boldsymbol{\mu}^{(i)} + \boldsymbol{\sigma}^{(i)} \odot \boldsymbol{\epsilon}^{(i,l)}$$

where

$$\boldsymbol{\epsilon}^{(i,l)} \sim \mathcal{N}(0, \mathbf{I})$$

4 Ex 4

The model summaries are given in figures 1 and 5 for latent dimensions of $k = 3$ and $k = 20$ respectively. For sampling from the generative network, a k dimensional vector \mathbf{z}_s was picked randomly in space \mathbb{R}^k . I looked at the mean and variance of the encoder distribution $q_\phi(\mathbf{z}|\mathbf{x})$ for the test data to decide on what would be reasonable values for the elements of \mathbf{z}_s . The 784 probability parameters of the Bernoulli distribution $p_\theta(\mathbf{x}|\mathbf{z})$ corresponding to this latent variable were obtained by using $\mathbf{p}_B = \text{decoder.predict}(\mathbf{z}_s)$. Then we simply independently sampled 784 random variables with these probabilities from the Bernoulli distribution i.e. the sample is given by $\mathbf{x}_s = \text{Bernoulli}(\mathbf{p}_B)$.

4.1 Number of latent variables K=3

- (a) Loss on test data: $KL \text{ loss} = 7.89$, Reconstruction error = 131.80
- (b) Some samples \mathbf{x}_s from the generative network for $k = 2$ are shown in figure 2
- (c) Some original test examples and their reconstructions are shown in figures 3 and 4.

4.2 Number of Latent variables $K = 20$

- (a) Loss on test data: $KL \text{ loss} = 25.23$, Reconstruction error = 78.33
- (b) Some samples \mathbf{x}_s from the generative network for $k = 20$ are shown in figure 6.
- (c) Some original test examples and their reconstructions are shown in figures 7 and 8.

We observe that as the dimension of the latent variable increases from $k = 3$

```
In [553]: vae.summary()
```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 2), (None, 2), (N 403972	
decoder (Model)	(None, 784)	403728
Total params: 807,700		
Trainable params: 807,700		
Non-trainable params: 0		

Figure 1




Figure 2: Samples \mathbf{x}_s for latent variable (a) $\mathbf{z}_s = [-2, 4]$ (b) $\mathbf{z}_s = [-0.2, 0.4]$ (c) $\mathbf{z}_s = [0.1, 0.1]$

to $k = 20$, the overall loss decreases. It is interesting to see that the KL loss actually increases with dimension so that the posterior $q_\phi(\mathbf{z}|\mathbf{x})$ is less close to the prior $p(\mathbf{z})$. However, the reconstruction error decreases with dimension.




Figure 3: (a) Original (b) Reconstructed




Figure 4: (a) Original (b) Reconstructed

```
In [80]: vae.summary()
```

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 20), (None, 20),	422440
decoder (Model)	(None, 784)	412944
Total params: 835,384		
Trainable params: 835,384		
Non-trainable params: 0		

Figure 5




Figure 6: Samples \boldsymbol{x}_s




Figure 7: (a) Original (b) Reconstructed




Figure 8: (a) Original (b) Reconstructed

Appendix

```
from keras import backend as K
from keras.losses import
    binary_crossentropy
from keras.datasets import mnist
from keras.layers import Lambda, Input,
    Dense
from keras.models import Model
import tensorflow as tf

import numpy as np
import matplotlib.pyplot as plt

#sampling function-reparametrization
    trick
def sampling(args):

    z_mean, z_log_var = args
    z_std = K.sqrt(K.exp(z_log_var))
    batch = K.shape(z_mean)[0]
    dim = K.int_shape(z_mean)[1]
    epsl = K.random_normal(shape=(batch,dim))
    return z_mean + z_std*epsl

#define losses
def reg_loss(args):

    z_mean, z_log_var = args
    #loss per datapoint
    kl_loss = 1 + z_log_var - K.exp(z_log_var)
        ) - K.square(z_mean)
    kl_loss = K.sum(kl_loss, axis = -1)
    kl_loss *= -0.5;
    #average over all datapoints
    kl_loss = K.mean(kl_loss)

    return kl_loss

def recn_loss(varargs):

    inputs, outputs = varargs
    #loss per data point assuming nos. of MC
        samples r=1
```

```

#recons_loss = K.sum(inputs * K.log(
    outputs) + (1-inputs)*K.log(1-outputs)
    , axis=-1)
recons_loss = binary_crossentropy(inputs ,
    outputs)
recons_loss*= original_dim
#average over all data points
recons_loss = K.mean(recons_loss)

return recons_loss

# MNIST dataset
(x_train , y_train), (x_test , y_test) =
    mnist.load_data()
image_size = x_train.shape[1]
original_dim = image_size * image_size
x_train = np.reshape(x_train , [-1,
    original_dim])
x_test = np.reshape(x_test , [-1,
    original_dim])
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# network parameters
input_shape = (original_dim , )
intermediate_dim = 512
batch_size = 128
latent_dim = 3
epochs = 20

# build encoder model
inputs = Input(shape=input_shape , name='
    encoder_input ')
x = Dense(intermediate_dim , activation='
    relu')(inputs)
z_mean = Dense(latent_dim , name='z_mean')
    (x)
z_log_var = Dense(latent_dim , name='
    z_log_var')(x)
z = Lambda(sampling , output_shape=(
    latent_dim ,), name='z')([z_mean ,
    z_log_var])
encoder = Model(inputs , [z_mean ,
    z_log_var , z] , name='encoder')

# build decoder model

```



```

latent_inputs = Input(shape=(latent_dim ,)
                        , name='z_sampling ')
x = Dense(intermediate_dim , activation='
relu ')(latent_inputs)
outputs = Dense(original_dim , activation
='sigmoid ')(x)
decoder = Model(latent_inputs , outputs ,
                name='decoder ')

outputs = decoder(encoder(inputs) [2])

vae = Model(inputs , outputs , name='vae ')
vae_loss = reg_loss ([z_mean , z_log_var ])
          + recn_loss ([inputs , outputs ])
vae.add_loss (vae_loss)
vae.compile(optimizer='adam' , metrics =
            ['reg_loss ' , 'recn_loss '])

#fit
history = vae.fit(x_train , epochs=epochs ,
                  batch_size=batch_size , validation_data
                  =(x_test , None))

#predict
ztest_mean , ztest_log_var , ztest =
    encoder.predict(x_test)
x_pred = decoder.predict(ztest)

#Ex4.1
x_pred = tf.convert_to_tensor(x_pred , np.
                              float32)
kl_loss = reg_loss ([ztest_mean ,
                    ztest_log_var ])
recons_loss = recn_loss ([x_test , x_pred])
with tf.Session() as sess:
    kl_loss = kl_loss.eval()

with tf.Session() as sess:
    recons_loss = recons_loss.eval()

print(kl_loss)
print(recons_loss)

#Ex4.2
def sample_generative_network(laten_var):

```

```

#probabilities of bernoulli distribution
prb = decoder.predict(laten_var)
#sample from bernoulli
xhat = np.random.binomial(1,prb)
xhat = np.reshape(xhat,(image_size ,
                    image_size))
plt.imshow(xhat)

#ex4.3
x_pred = decoder.predict(ztest);
idx = 700
#original image
plt.imshow(np.reshape(x_test[idx,:],(
    image_size , image_size)))
plt.show()
#reconstructed image
plt.show()
plt.imshow(np.reshape(x_pred[idx,:],(
    image_size , image_size)))

```